

Shivansh Kalra

+91-9717094960 | linkedin.com/in/shivansh-kalra | github.com/shivansh193

TECHNICAL SKILLS

Languages: JavaScript/TypeScript, Python, Node.js, C++, SQL, Bash

Full-Stack: React, Next.js, Electron (learning), Express.js, NestJS, WebSockets, REST APIs

AI/ML: PyTorch, TensorFlow, OpenAI/Anthropic APIs, LangChain, Vector DBs, RAG Systems

Infrastructure: AWS (EC2, S3, Lambda), Docker, Linux, Redis, PostgreSQL, MongoDB, CI/CD

Core Strengths: Fast prototyping, system design, 0-1 products, end-to-end ownership

PROOF OF WORK: FAST SHIPPING & TECHNICAL OWNERSHIP

• InVariant

Founder & Solo Developer

In Development

Nov. 2024 – Present

- Building infrastructure and security analysis tool for AI-generated code from scratch, owning everything from system design to UI/UX to deployment pipeline.
- Implementing Python-based static analysis engine that processes code repositories, identifies vulnerabilities, and generates actionable security reports with 90%+ accuracy.
- Shipping fast iteration cycles (weekly releases), gathering user feedback from developers, and pivoting features based on real-world usage patterns.
- Strong opinions on AI tooling UX: most tools prioritize "impressive" over "useful." Building for what developers actually need, not what looks good in demos.

• FilFlo

Backend Engineer (Microservices Migration)

Remote

Sept. 2024 – Present

- Architected and migrated 50% of monolithic codebase to microservices using Node.js/NestJS, reducing deployment time by 40% and enabling independent service scaling.
- Optimized core systems: implemented Redis caching layer achieving 37% performance improvement, reduced API response times from 800ms to <300ms for 10K+ daily requests.
- Built production APIs with comprehensive error handling, monitoring, and logging, maintaining 99.8% uptime across distributed services.
- Owned end-to-end features: from database schema design to API implementation to Docker deployment, shipping autonomously with minimal oversight.

• Testline

AI Product Engineer

Remote

Dec. 2024 – Feb. 2025

- Shipped 3 AI-powered features serving 3,000+ users in 8 weeks, integrating OpenAI GPT-4 and Google Gemini APIs with custom prompt engineering and context management.
- Built LLM-powered agent using ReAct framework for educational tutoring, handling complex multi-turn conversations and maintaining context across sessions.
- Implemented efficient data pipelines for processing user inputs, managing token limits, and caching responses to optimize cost and latency.
- Wrote production-grade code with 85% test coverage, identifying and fixing 25+ edge cases before launch through rigorous testing.

• Frontrunner

Full-Stack Engineer (Rapid Prototyping)

Remote

May 2024 – Aug. 2024

- Prototyped and launched 2 full-stack MVPs in 8-week sprints using React, Node.js, and MongoDB, contributing to 40% revenue increase through fast feature delivery.
- Shipped end-to-end: designed schemas, built APIs, implemented frontends, deployed to production, and iterated based on user feedback, all autonomously.
- Built automated CI/CD pipelines reducing deployment overhead by 70%, enabling multiple daily releases without breaking production.

• CatalystIQ

Founder & Technical Lead (0 to Revenue in 3 Months)

Gurgaon, India

Jan. 2024 – April 2024

- Built AI-powered SaaS platform from scratch, shipping enterprise-ready product that generated 1 Lakh+ revenue with 10+ B2B clients within first quarter.
- Owned full stack: React frontend, Node.js backend, Python ML services, PostgreSQL database, AWS infrastructure, achieving 99%+ uptime with zero DevOps support.
- Implemented ML-based recruitment matching system using PyTorch, reducing hiring cycle time by 65% through custom neural architecture and data pipeline optimization.
- Opinionated on product-market fit: killed 3 features after user research showed low adoption, doubled down on what users actually paid for.

• Petner	Remote
<i>Founder & Solo Developer (0 to 25 Partnerships)</i>	<i>Apr. 2023 – Dec. 2023</i>
– Built and shipped full B2B SaaS platform solo, securing 25+ partnerships and onboarding 100+ users through product-led growth strategy.	
– Implemented offline-first architecture with real-time sync using Firebase, ensuring reliability even with poor network conditions.	
– Rapid experimentation: shipped 15+ feature iterations based on user feedback, maintaining 2-week release cycles throughout 8-month product lifecycle.	

NOTABLE TECHNICAL PROJECTS

- **SNUXplore** | *RAG-based AI platform serving 30K+ concurrent users*
 - Built scalable AI platform with RAG architecture (Azure OpenAI + vector embeddings), handling 30K+ concurrent users with sub-100ms query response times.
 - Designed distributed system with load balancing, Redis caching, and optimized database queries, proving ability to build production-scale infrastructure.
 - Strong opinions on context management: most RAG systems are overkill. Shipped lean retrieval pipeline that achieved 92% answer relevance with minimal complexity.
- **BuildFlow** | *SaaS platform with 200+ enterprise users, 500+ projects*
 - Shipped production SaaS in 6 weeks with full RBAC, multi-tenant architecture, and real-time collaboration features, demonstrating fast execution with quality.
 - Built automated CI/CD pipeline maintaining 99.9% uptime, handling deployments without downtime using blue-green deployment strategy.
- **AI Agent Experiments** | *Personal research on LLM-powered autonomous systems*
 - Built experimental desktop automation agent using Python, exploring challenges in reliable computer use, context management, and user intent interpretation.
 - Prototyped memory systems for long-term context retention across sessions, testing vector databases and compression techniques for efficient storage/retrieval.
 - Strong opinions on agent UX: most demos fail because they prioritize autonomy over predictability. Users want agents that work WITH them, not FOR them.

EDUCATION & BACKGROUND

• Shiv Nadar University	Noida, India
<i>B.Tech Computer Science, Expected Graduation: May 2027</i>	<i>Aug. 2023 – Present</i>
– 2+ years of hands-on building experience across startups, internships, and personal projects. Self-taught most of my stack through shipping real products.	

WHY I'M A GOOD FIT

Fast shipper: Built 6+ products from 0 to users/revenue, comfortable with ambiguity and high velocity

Full ownership: Don't need hand-holding. Give me a problem, I'll design, build, deploy, and iterate until it works

Opinionated & curious: Strong views on AI UX, system design, and what makes products actually useful (not just impressive)

Proof over titles: 0-2 YOE formally, but 2+ years of real shipping experience. Judge me by what I've built, not credentials

Remote-first mindset: All my work has been remote, async, high-trust environments. Thrive with autonomy and clear goals